

Sentinel Agentic AML Investigation Case Manager

UiPath AgentHack Track 1: Maestro Case. Build the AI agents of tomorrow.



Team: TheSecond
Project Title: Sentinel



ZEMBA Wendemi Brice Roméo

Assitant Researcher @ University Mohamed VI Polytechnic
Wendemi.ZEMBA-EXT@um6p.ma

Problem statement and proposed solution


Problem

- ~95% of AML transaction-monitoring alerts are false positives.
- Each one is still investigated by hand — pull data, screen sanctions, write the SAR.
- A compliance officer must sign off every filing — it's legally mandated.
- Result: slow, inconsistent, costly, and the risk of late or defective filings.

Solution

- Sentinel makes every alert a governed case in UiPath Maestro.
- Agents investigate, robots act on systems of record, an MLRO signs off.
- The path is dynamic: duplicates auto-close, sanctions fast-track, weak cases loop back.
- Every decision is auditable end to end — the differentiator for regulators.

How it works: one case, five stages



Dynamic by design: duplicates auto-close · sanctions fast-track to a senior MLRO · Quality Review loops back to Investigation · SLA breaches escalate.

Solution architecture

UiPath Maestro Case orchestration · SLAs · escalation · full audit trail

AGENTS

- **Triage (low-code, Agent Builder)**
- Investigator (coded · LangGraph)
- Narrator (coded · LangGraph)
- Quality Review (low-code)

ROBOTS

- **DataFetcher (API Workflow)**
- Filer (RPA — files the SAR)
- Notifier (escalations)
- DocIntake (Document Understanding)

PEOPLE

- **MLRO / Compliance Officer**
- Reviews in Action Center
- Owns the filing decision
- Accountable · audited · in control

Agents never touch source systems — robots do, behind the UiPath credential vault. A CrewAI build of the Investigator drops in behind the same Maestro stage.

Benefits and technologies used

Details

End-user	AML analysts and the MLRO / compliance officer
User department	Financial Crime / BSA-AML Compliance
Industries	Banking, fintech, payments, crypto, gaming
UiPath products used	Maestro Case, Agent Builder, Coded Agents (Python SDK), Action Center, API Workflows + RPA, Document Understanding, Orchestrator
Other — APIs / tech	LangGraph, CrewAI, Claude (Anthropic); built via Claude Code + uip CLI

Benefits, impact and outcomes

- Multi-hour manual investigations → minutes of agent work + one human decision.
- Consistent rigor on every alert — no analyst-to-analyst drift.
- Reconstructable risk score + complete audit trail → fewer late / defective filings.
- A human is accountable for every SAR filing and dismissal.
- Production path: swap mocks for core-banking, sanctions, and FinCEN E-Filing.

Platform depth and coding-agent bonus

Hybrid agents, one control plane

Low-code Agent Builder (Triage, QA) + coded LangGraph agents (Investigator, Narrator) + RPA robots + Document Understanding — all governed by Maestro Case.

External frameworks, governed

The Investigator ships in two builds — LangGraph and CrewAI — behind the same UiPath contract, swappable at the Maestro stage. UiPath stays the orchestration layer.

Built with coding agents (+bonus)

The coded agents, risk logic, tests, and deploy scripts were built with Claude Code via the UiPath uip CLI. Documented and verifiable in CODING_AGENTS.md.

Engineered to run, not just demo

Both coded agents pass automated tests (8 total) and run offline with a deterministic fallback — a case is never left without a result.

THANK YOU.

Sentinel agents investigate, robots act, a compliance officer signs off.

GitHub: <https://github.com/BriceZemba/sentinel-aml> · Demo video & deck in the submission.